

# DAT 1


26/3/16

# Agenda

- Learning Theory
- Reinforcement Learning Intro
- Spark
- (Recommender Practice)
- Random Topics

# Learning Theory

Bias / Variance: parameters, models



How much data needed?

Which models to even consider?

# Learning Theory

**Lemma.** (The union bound). Let  $A_1, A_2, \dots, A_k$  be  $k$  different events (that may not be independent). Then

$$P(A_1 \cup \dots \cup A_k) \leq P(A_1) + \dots + P(A_k).$$

**Lemma.** (Hoeffding inequality) Let  $Z_1, \dots, Z_m$  be  $m$  independent and identically distributed (iid) random variables drawn from a Bernoulli( $\phi$ ) distribution. I.e.,  $P(Z_i = 1) = \phi$ , and  $P(Z_i = 0) = 1 - \phi$ . Let  $\hat{\phi} = (1/m) \sum_{i=1}^m Z_i$  be the mean of these random variables, and let any  $\gamma > 0$  be fixed. Then

$$P(|\phi - \hat{\phi}| > \gamma) \leq 2 \exp(-2\gamma^2 m)$$

How many training points (samples) to be within 0.01 of each other with probability 95%?

# Learning Theory

Binary classification empirical error

$$\hat{\varepsilon}(h) = \frac{1}{m} \sum_{i=1}^m 1\{h(x^{(i)}) \neq y^{(i)}\}.$$

(fraction misclassified)

Generalisation error

$$\varepsilon(h) = P_{(x,y) \sim \mathcal{D}}(h(x) \neq y).$$

# Learning Theory

PAC framework  
“probably approximately correct”

i.e. training and testing from same distribution

# Learning Theory

Minimising training error

$$\hat{\theta} = \arg \min_{\theta} \hat{e}(h_{\theta}).$$

$$\mathcal{H} = \{h_{\theta} : h_{\theta}(x) = 1\{\theta^T x \geq 0\}, \theta \in \mathbb{R}^{n+1}\}$$

equivalent to  $\hat{h} = \arg \min_{h \in \mathcal{H}} \hat{e}(h)$

# Learning Theory

Finite  $\mathcal{H} = \{h_1, \dots, h_k\}$

We want guarantees on generalisation error when we minimise empirical error

- theoretical framework
- applied cases e.g. too few data points for cross-validation



# Learning Theory

- Take fixed  $h_i$  in  $H$
- Let  $Z_j = 1\{h_i(x^{(j)}) \neq y^{(j)}\}$ .
- Training error of  $h_i$

$$\hat{\varepsilon}(h_i) = \frac{1}{m} \sum_{j=1}^m Z_j.$$

- Hoeffding

$$P(|\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) \leq 2 \exp(-2\gamma^2 m).$$

# Learning Theory

Great. But we don't just want this for one  $h_i$ , we want it for all of  $H$  simultaneously

$$P(A_i) \quad \dots \quad P(|\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) \leq 2 \exp(-2\gamma^2 m).$$

$$\begin{aligned} P(\exists h \in \mathcal{H}. |\varepsilon(h) - \hat{\varepsilon}(h)| > \gamma) &= P(A_1 \cup \dots \cup A_k) \\ &\leq \sum_{i=1}^k P(A_i) \\ &\leq \sum_{i=1}^k 2 \exp(-2\gamma^2 m) \\ &= 2k \exp(-2\gamma^2 m) \end{aligned}$$

# Learning Theory

$$\begin{aligned}P(\exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(A_1 \cup \dots \cup A_k) \\&\leq \sum_{i=1}^k P(A_i) \\&\leq \sum_{i=1}^k 2 \exp(-2\gamma^2 m) \\&= 2k \exp(-2\gamma^2 m)\end{aligned}$$

If we subtract both sides from 1, we find that

$$\begin{aligned}P(\neg \exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(\forall h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| \leq \gamma) \\&\geq 1 - 2k \exp(-2\gamma^2 m)\end{aligned}$$

# Learning Theory

## Example

For instance, we can ask the following question: Given  $\gamma$  and some  $\delta > 0$ , how large must  $m$  be before we can guarantee that with probability at least  $1 - \delta$ , training error will be within  $\gamma$  of generalization error? By setting  $\delta = 2k \exp(-2\gamma^2 m)$  and solving for  $m$ , [you should convince yourself this is the right thing to do!], we find that if

$$m \geq \frac{1}{2\gamma^2} \log \frac{2k}{\delta},$$

then with probability at least  $1 - \delta$ , we have that  $|\varepsilon(h) - \hat{\varepsilon}(h)| \leq \gamma$  for all  $h \in \mathcal{H}$ . (Equivalently, this shows that the probability that  $|\varepsilon(h) - \hat{\varepsilon}(h)| > \gamma$  for some  $h \in \mathcal{H}$  is at most  $\delta$ .) This bound tells us how many training examples we need in order make a guarantee. The training set size  $m$  that a certain method or algorithm requires in order to achieve a certain level of performance is also called the algorithm's **sample complexity**.

# Learning Theory

With just a bit more effort, we can get an explicit guarantee on generalisation error

**Theorem.** Let  $|\mathcal{H}| = k$ , and let any  $m, \delta$  be fixed. Then with probability at least  $1 - \delta$ , we have that

$$\varepsilon(\hat{h}) \leq \left( \min_{h \in \mathcal{H}} \varepsilon(h) \right) + 2\sqrt{\frac{1}{2m} \log \frac{2k}{\delta}}.$$

lower bias: increasing  
H class

variance tradeoff of  
increasing H class

# Learning Theory

In summary, given a hypothesis class of complexity  $k$ , we can figure out how many training examples we need to be within a certain distance from minimising generalisation error when carrying out empirical risk minimisation

Analog for continuous case:  
Vapnik-Chervonenki (VC) dimension

# Reinforcement Learning

Previously,  $y$  had a “correct” value

Now instead, what if the best we can provide is a reward function

e.g. robot motion, gaming...trading?

<https://www.youtube.com/watch?v=qv6UVQQ0F44>

<https://www.youtube.com/watch?v=M8YjvHYbZ9w>

<https://www.youtube.com/watch?v=rVlhMGQgDkY>

# Reinforcement Learning

A Markov decision process is a tuple  $(S, A, \{P_{sa}\}, \gamma, R)$ , where:

- $S$  is a set of **states**. (For example, in autonomous helicopter flight,  $S$  might be the set of all possible positions and orientations of the helicopter.)
- $A$  is a set of **actions**. (For example, the set of all possible directions in which you can push the helicopter's control sticks.)
- $P_{sa}$  are the state transition probabilities. For each state  $s \in S$  and action  $a \in A$ ,  $P_{sa}$  is a distribution over the state space. We'll say more about this later, but briefly,  $P_{sa}$  gives the distribution over what states we will transition to if we take action  $a$  in state  $s$ .
- $\gamma \in [0, 1)$  is called the **discount factor**.
- $R : S \times A \mapsto \mathbb{R}$  is the **reward function**. (Rewards are sometimes also written as a function of a state  $S$  only, in which case we would have  $R : S \mapsto \mathbb{R}$ ).



# Reinforcement Learning

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$
$$s_1 \sim P_{s_0 a_0}.$$

Upon visiting the sequence of states  $s_0, s_1, \dots$  with actions  $a_0, a_1, \dots$ , our total payoff is given by

$$R(s_0, a_0) + \gamma R(s_1, a_1) + \gamma^2 R(s_2, a_2) + \dots .$$

Or, when we are writing rewards as a function of the states only, this becomes

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots .$$

# Reinforcement Learning

We want to maximise expected discounted reward

Economic interpretation of gamma?

$$E [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

# Reinforcement Learning

**Policy** to choose actions

$$\pi : S \mapsto A$$

$$a = \pi(s)$$

**Value function** based on policy

$$V^\pi(s) = \mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots \mid s_0 = s, \pi].$$

# Reinforcement Learning

$$V^\pi(s) = \mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots \mid s_0 = s, \pi].$$

## **Bellman equation**

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s').$$

Can you draw this (perhaps as a decision tree?)

# Reinforcement Learning

We want to find the optimal value

$$V^*(s) = \max_{\pi} V^{\pi}(s).$$

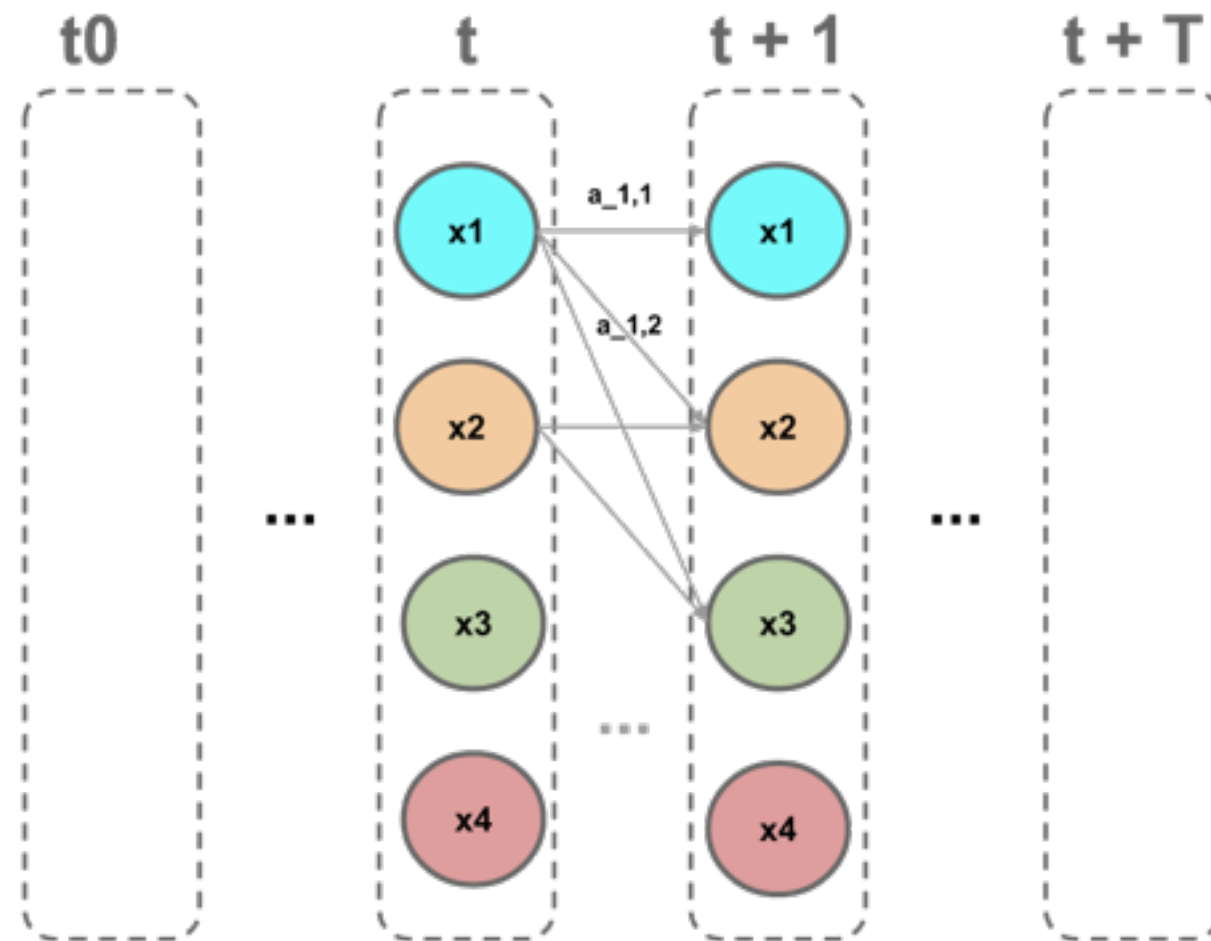
or

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s').$$

max over possible actions

# Reinforcement Learning

Exponential paths to explore



Hansel and Grettel!

# Reinforcement Learning

Recommended for Dynamic Programming and  
Value Iteration:

<https://en.wikipedia.org/wiki/Memoization>

<https://www.youtube.com/watch?v=oefOCk3koZo>

<https://www.youtube.com/watch?v=ip4iSMRW5X4>

# Reinforcement Learning

Extensions:

Learning  $P_{sa}$

Continuous space (discretisation)



**SPARK**

# Random Topics


## *Neural Networks*

A computational system  
comprised of layers and each layer  
is built of interconnected perceptrons

Built to model the animal nervous system?



# Random Topics

## *Neural Networks*



# Random Topics

## *Neural Networks*




# Random Topics

## *Neural Networks*

### Single Perceptron


$$f_{log}(z) = \frac{1}{1 + e^{-z}}$$

$f_{log}$  is called **logistic function**



Takes in input and uses an activation function in order to output


# Random Topics *Neural Networks*



# Random Topics *Neural Networks*

Artificial Neural Networks are  
also known as multi layer perceptrons

A **multi layer perceptrons (MLP)** is a finite acyclic graph. The nodes are neurons with logistic activation.



Random Topics  
*Neural Networks*

**But how does it learn?!**

**Back-Propagation**

<https://en.wikipedia.org/wiki/Backpropagation>



# Random Topics

## *Neural Networks*

### Pros

- Online model (updates as you go)
  - Doesn't need to be fit all of the time
- Very fast predictions
- Can approximate almost any type of function
- Can be used in a supervised and unsupervised manner
- Super cool

### Cons

- Requires many training samples to be considered good
- Hard to describe what is happening
- Requires a lot of hardware / computation power
- Slow to train
- Can be difficult to use

# Random Topics

## *Neural Networks*

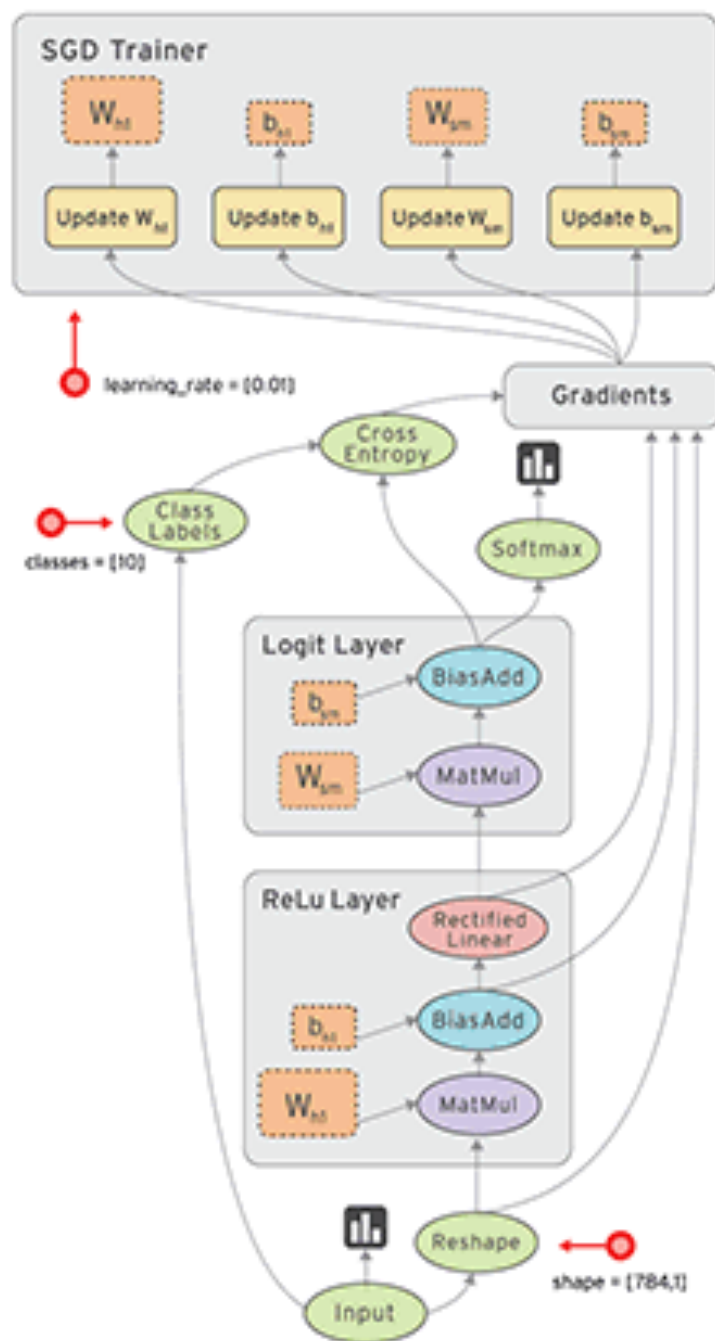
<https://www.tensorflow.org/>

TensorFlow is an Open Source Software  
Library for Machine Intelligence

GET STARTED

# Random Topics

## *Neural Networks*



Random Topics  
*Web Scraping*

Python BeautifulSoup

<http://www.crummy.com/software/BeautifulSoup/bs4/doc/>

Exercise: scrape this

list of buyers

<http://econpy.pythonanywhere.com/ex/001.html>

finance

<http://stackoverflow.com/questions/10040954/alternative-to-google-finance-api>

# Random Topics

## *Web Scraping*

### Concepts

- Robustness: changes in the web page/api
- Scale: massive web scraping fault tolerance
- Integrity: external data sources

# Random Topics

## *Setting up servers*





1. AWS account
2. ec2 instance
3. ssh into it



1. heroku account
2. python "hello" api
3. push to heroku

# Random Topics

## ***Resources Going Forward***



# Random Topics

## ***Resources Going Forward***

Michael Nielsen, Christopher Olah  
<http://colah.github.io/>



Geoffrey Hinton



Yann LeCun



Random Topics  
***Resources Going Forward***



Andrej Karpathy



Hugo Larochelle

Random Topics  
***Resources Going Forward***

Linear Algebra

Probability

Parallel Programming

Computer Science

Topology

Information Theory